

TOP 50 INTERVIEW QUESTIONS FOR SQL

1. What is SQL? How is it different from MySQL or PostgreSQL?
2. What are the different types of SQL statements?
3. Explain the difference between WHERE and HAVING.
4. What are PRIMARY KEY, FOREIGN KEY, UNIQUE, and CHECK constraints?
5. What is the difference between DELETE, TRUNCATE, and DROP?
6. What is normalization? Explain different normal forms.
7. What is denormalization and when is it useful?
8. Explain the difference between CHAR and VARCHAR.
9. What are ACID properties in databases?
10. What is the difference between INNER JOIN, LEFT JOIN, RIGHT JOIN, and FULL JOIN?
11. Write a query to find the second highest salary from an Employee table.
12. Write a query to get the department-wise average salary.
13. How would you retrieve duplicate records from a table?
14. How do you update a column with a calculation (e.g., 10% tax added)?
15. How would you delete only duplicate rows from a table?
16. Write a query to list customers who have placed more than 5 orders.
17. Write a query to join three or more tables.
18. What is a subquery? How is it different from a JOIN?
19. What is a correlated subquery? Give an example.
20. How do you filter data based on a date range?
21. What are WINDOW FUNCTIONS? Name a few.
22. What is the use of RANK(), DENSE_RANK(), and ROW_NUMBER()? 23. What is a Common Table Expression (CTE)? How is it different from a subquery? 24. What are stored procedures? When should they be used? 25. What is a trigger? Give a real-world example. 26. What is a VIEW? What are its pros and cons? 27. What are indexes? How do they improve performance? 28. What is a materialized view? 29. What are transactions? Explain COMMIT, ROLLBACK, and SAVEPOINT. 30. What are aggregate functions? List a few with examples.
31. How can you optimize a slow-running SQL query?
32. What is the EXPLAIN or EXPLAIN PLAN statement used for?
33. How does indexing affect INSERT, UPDATE, and DELETE performance?
34. What is a composite index and when should it be used?
35. What is normalization overhead and how do you deal with it?
36. How do you avoid Cartesian products in JOINS?
37. What is partitioning in SQL?
38. What causes a deadlock in SQL, and how can you prevent it?
39. What is the difference between clustered and non-clustered indexes?
40. What tools do you use to monitor SQL query performance?
41. You are asked to design a student-course grading system. What tables and relationships would you create?
42. How would you store and retrieve attendance for employees in a scalable way?
43. In a library system, how would you track overdue books and fines using SQL?
44. What would you do if a production database is missing some records due to a failed update?
45. How would you implement role-based access to sensitive information in SQL?
46. You are given a raw CSV with dirty data. How would you load and clean it using SQL?
47. How would you calculate monthly retention from a user login dataset?
48. What measures would you take to secure a database with sensitive data?
49. How do you create daily backup and restore plans for a SQL database?

